

Perform the following operations using 'python' Language on the Air quality dataset.

1. Data cleaning
2. Data integration
3. Data transformation
4. Data model Building

```
In [1]: import io
import pandas as pd
import numpy as np
import seaborn as sns
from scipy import stats
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier

rawdata_df = pd.read_csv(r"C:\Users\yasha\Downloads\AirQualityUCI.csv")
rawdata_df
```

Out[1]:

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)
0	3/10/2004	18:00:00	2.6	1360.0	150.0	11.9	1046.0	166.0
1	3/10/2004	19:00:00	2.0	1292.0	112.0	9.4	955.0	103.0
2	3/10/2004	20:00:00	2.2	1402.0	88.0	9.0	939.0	131.0
3	3/10/2004	21:00:00	2.2	1376.0	80.0	9.2	948.0	172.0
4	3/10/2004	22:00:00	1.6	1272.0	51.0	6.5	836.0	131.0
...
9466	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9467	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9468	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9469	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9470	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

9471 rows × 17 columns



```
In [2]: rawdata_df.isnull().sum()
```

```
Out[2]: Date          114  
Time            114  
CO(GT)          114  
PT08.S1(CO)     114  
NMHC(GT)        114  
C6H6(GT)        114  
PT08.S2(NMHC)   114  
NOx(GT)         114  
PT08.S3(NOx)    114  
NO2(GT)         114  
PT08.S4(NO2)    114  
PT08.S5(O3)     114  
T              114  
RH             114  
AH             114  
Unnamed: 15     9471  
Unnamed: 16     9471  
dtype: int64
```

```
In [3]: rawdata_df.columns
```

```
Out[3]: Index(['Date', 'Time', 'CO(GT)', 'PT08.S1(CO)', 'NMHC(GT)', 'C6H6(GT)',  
              'PT08.S2(NMHC)', 'NOx(GT)', 'PT08.S3(NOx)', 'NO2(GT)', 'PT08.S4(NO  
2)',  
              'PT08.S5(O3)', 'T', 'RH', 'AH', 'Unnamed: 15', 'Unnamed: 16'],  
              dtype='object')
```

```
In [108]: # data cleaning
```

```
In [4]: #we seleted only the columns having values
selected_columns=['Date', 'Time', 'CO(GT)', 'PT08.S1(CO)', 'NMHC(GT)', 'C6H6(GT)']
analysis_df = rawdata_df[selected_columns].copy()
analysis_df
```

Out[4]:

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)
0	3/10/2004	18:00:00	2.6	1360.0	150.0	11.9	1046.0	166.0
1	3/10/2004	19:00:00	2.0	1292.0	112.0	9.4	955.0	103.0
2	3/10/2004	20:00:00	2.2	1402.0	88.0	9.0	939.0	131.0
3	3/10/2004	21:00:00	2.2	1376.0	80.0	9.2	948.0	172.0
4	3/10/2004	22:00:00	1.6	1272.0	51.0	6.5	836.0	131.0
...
9466	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9467	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9468	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9469	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9470	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

9471 rows × 15 columns

```
In [5]: analysis_df.shape
```

Out[5]: (9471, 15)

In [6]: `analysis_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9471 entries, 0 to 9470
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                   9357 non-null   object
1   Time                   9357 non-null   object
2   CO(GT)                 9357 non-null   float64
3   PT08.S1(CO)            9357 non-null   float64
4   NMHC(GT)               9357 non-null   float64
5   C6H6(GT)               9357 non-null   float64
6   PT08.S2(NMHC)          9357 non-null   float64
7   NOx(GT)                9357 non-null   float64
8   PT08.S3(NOx)           9357 non-null   float64
9   NO2(GT)                9357 non-null   float64
10  PT08.S4(NO2)           9357 non-null   float64
11  PT08.S5(O3)            9357 non-null   float64
12  T                       9357 non-null   float64
13  RH                      9357 non-null   float64
14  AH                      9357 non-null   float64
dtypes: float64(13), object(2)
memory usage: 1.1+ MB
```

In [7]: `#dropping the all null values`
`analysis_df=analysis_df.dropna()`
`analysis_df`

Out[7]:

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)
0	3/10/2004	18:00:00	2.6	1360.0	150.0	11.9	1046.0	166.0
1	3/10/2004	19:00:00	2.0	1292.0	112.0	9.4	955.0	103.0
2	3/10/2004	20:00:00	2.2	1402.0	88.0	9.0	939.0	131.0
3	3/10/2004	21:00:00	2.2	1376.0	80.0	9.2	948.0	172.0
4	3/10/2004	22:00:00	1.6	1272.0	51.0	6.5	836.0	131.0
...
9352	4/4/2005	10:00:00	3.1	1314.0	-200.0	13.5	1101.0	472.0
9353	4/4/2005	11:00:00	2.4	1162.0	-200.0	11.4	1027.0	353.0
9354	4/4/2005	12:00:00	2.4	1142.0	-200.0	12.4	1062.0	293.0
9355	4/4/2005	13:00:00	2.1	1002.0	-200.0	9.5	960.0	234.0
9356	4/4/2005	14:00:00	2.2	1071.0	-200.0	11.9	1047.0	265.0

9357 rows × 15 columns



```
In [113]: analysis_df.isnull().sum()
```

```
Out[113]: Date                0
          Time                0
          CO(GT)              0
          PT08.S1(CO)         0
          NMHC(GT)            0
          C6H6(GT)            0
          PT08.S2(NMHC)       0
          NOx(GT)             0
          PT08.S3(NOx)        0
          NO2(GT)             0
          PT08.S4(NO2)        0
          PT08.S5(O3)         0
          T                   0
          RH                  0
          AH                  0
          dtype: int64
```

```
In [114]: #data transformation
          analysis_df['Date']
```

```
Out[114]: 0      3/10/2004
          1      3/10/2004
          2      3/10/2004
          3      3/10/2004
          4      3/10/2004
          ...
          9352   4/4/2005
          9353   4/4/2005
          9354   4/4/2005
          9355   4/4/2005
          9356   4/4/2005
          Name: Date, Length: 9357, dtype: object
```

```
In [115]: analysis_df['Date'] = pd.to_datetime(analysis_df.Date)
analysis_df['Date']
```

C:\Users\yasha\AppData\Local\Temp\ipykernel_3164\2240121029.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
analysis_df['Date'] = pd.to_datetime(analysis_df.Date)
```

```
Out[115]: 0      2004-03-10
1      2004-03-10
2      2004-03-10
3      2004-03-10
4      2004-03-10
...
9352   2005-04-04
9353   2005-04-04
9354   2005-04-04
9355   2005-04-04
9356   2005-04-04
Name: Date, Length: 9357, dtype: datetime64[ns]
```

```
In [116]: analysis_df['Year'] = pd.DatetimeIndex(analysis_df['Date']).year
analysis_df['month'] = pd.DatetimeIndex(analysis_df['Date']).month

analysis_df
```

C:\Users\yasha\AppData\Local\Temp\ipykernel_3164\1901626825.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
analysis_df['Year'] = pd.DatetimeIndex(analysis_df['Date']).year
```

C:\Users\yasha\AppData\Local\Temp\ipykernel_3164\1901626825.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
analysis_df['month'] = pd.DatetimeIndex(analysis_df['Date']).month
```

Out[116]:

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	F
0	2004-03-10	18:00:00	2.6	1360.0	150.0	11.9	1046.0	166.0	
1	2004-03-10	19:00:00	2.0	1292.0	112.0	9.4	955.0	103.0	
2	2004-03-10	20:00:00	2.2	1402.0	88.0	9.0	939.0	131.0	
3	2004-03-10	21:00:00	2.2	1376.0	80.0	9.2	948.0	172.0	
4	2004-03-10	22:00:00	1.6	1272.0	51.0	6.5	836.0	131.0	
...	
9352	2005-04-04	10:00:00	3.1	1314.0	-200.0	13.5	1101.0	472.0	
9353	2005-04-04	11:00:00	2.4	1162.0	-200.0	11.4	1027.0	353.0	
9354	2005-04-04	12:00:00	2.4	1142.0	-200.0	12.4	1062.0	293.0	
9355	2005-04-04	13:00:00	2.1	1002.0	-200.0	9.5	960.0	234.0	
9356	2005-04-04	14:00:00	2.2	1071.0	-200.0	11.9	1047.0	265.0	

9357 rows × 17 columns



```
In [117]: analysis_df['month_apha'] = pd.to_datetime(analysis_df['month'], format='%m')
analysis_df
```

C:\Users\yasha\AppData\Local\Temp\ipykernel_3164\675708319.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
analysis_df['month_apha'] = pd.to_datetime(analysis_df['month'], format='%m').dt.month_name()
```

```
Out[117]:
```

	Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	F
0	2004-03-10	18:00:00	2.6	1360.0	150.0	11.9	1046.0	166.0	
1	2004-03-10	19:00:00	2.0	1292.0	112.0	9.4	955.0	103.0	
2	2004-03-10	20:00:00	2.2	1402.0	88.0	9.0	939.0	131.0	
3	2004-03-10	21:00:00	2.2	1376.0	80.0	9.2	948.0	172.0	
4	2004-03-10	22:00:00	1.6	1272.0	51.0	6.5	836.0	131.0	
...	
9352	2005-04-04	10:00:00	3.1	1314.0	-200.0	13.5	1101.0	472.0	
9353	2005-04-04	11:00:00	2.4	1162.0	-200.0	11.4	1027.0	353.0	
9354	2005-04-04	12:00:00	2.4	1142.0	-200.0	12.4	1062.0	293.0	
9355	2005-04-04	13:00:00	2.1	1002.0	-200.0	9.5	960.0	234.0	
9356	2005-04-04	14:00:00	2.2	1071.0	-200.0	11.9	1047.0	265.0	

9357 rows × 18 columns




```
In [118]: #integrate
data1 = analysis_df[['Date', 'Time', 'CO(GT)']].loc[0:15]
data1
```

```
Out[118]:
```

	Date	Time	CO(GT)
0	2004-03-10	18:00:00	2.6
1	2004-03-10	19:00:00	2.0
2	2004-03-10	20:00:00	2.2
3	2004-03-10	21:00:00	2.2
4	2004-03-10	22:00:00	1.6
5	2004-03-10	23:00:00	1.2
6	2004-03-11	0:00:00	1.2
7	2004-03-11	1:00:00	1.0
8	2004-03-11	2:00:00	0.9
9	2004-03-11	3:00:00	0.6
10	2004-03-11	4:00:00	-200.0
11	2004-03-11	5:00:00	0.7
12	2004-03-11	6:00:00	0.7
13	2004-03-11	7:00:00	1.1
14	2004-03-11	8:00:00	2.0
15	2004-03-11	9:00:00	2.2

```
In [119]: data2 = analysis_df[['Date', 'Time', 'CO(GT)']].loc[16:30]
data2
```

```
Out[119]:
```

	Date	Time	CO(GT)
16	2004-03-11	10:00:00	1.7
17	2004-03-11	11:00:00	1.5
18	2004-03-11	12:00:00	1.6
19	2004-03-11	13:00:00	1.9
20	2004-03-11	14:00:00	2.9
21	2004-03-11	15:00:00	2.2
22	2004-03-11	16:00:00	2.2
23	2004-03-11	17:00:00	2.9
24	2004-03-11	18:00:00	4.8
25	2004-03-11	19:00:00	6.9
26	2004-03-11	20:00:00	6.1
27	2004-03-11	21:00:00	3.9
28	2004-03-11	22:00:00	1.5
29	2004-03-11	23:00:00	1.0
30	2004-03-12	0:00:00	1.7

```
In [120]: integrate=pd.merge(data1,data2,on='Date',how='inner')
integrate
```

```
Out[120]:
```

	Date	Time_x	CO(GT)_x	Time_y	CO(GT)_y
0	2004-03-11	0:00:00	1.2	10:00:00	1.7
1	2004-03-11	0:00:00	1.2	11:00:00	1.5
2	2004-03-11	0:00:00	1.2	12:00:00	1.6
3	2004-03-11	0:00:00	1.2	13:00:00	1.9
4	2004-03-11	0:00:00	1.2	14:00:00	2.9
...
135	2004-03-11	9:00:00	2.2	19:00:00	6.9
136	2004-03-11	9:00:00	2.2	20:00:00	6.1
137	2004-03-11	9:00:00	2.2	21:00:00	3.9
138	2004-03-11	9:00:00	2.2	22:00:00	1.5
139	2004-03-11	9:00:00	2.2	23:00:00	1.0

140 rows × 5 columns

```
In [23]: analysis_df['AH'] = analysis_df['AH'].apply(lambda x:0 if x<=0 else 1)
analysis_df
```

C:\Users\yasha\AppData\Local\Temp\ipykernel_6980\768665957.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
analysis_df['AH'] = analysis_df['AH'].apply(lambda x:0 if x<=0 else 1)
```

Out[23]:

Date	Time	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	PT08.S1(NMHC)
3/10/2004	18:00:00	2.6	1360.0	150.0	11.9	1046.0	166.0	104.0
3/10/2004	19:00:00	2.0	1292.0	112.0	9.4	955.0	103.0	103.0
3/10/2004	20:00:00	2.2	1402.0	88.0	9.0	939.0	131.0	131.0
3/10/2004	21:00:00	2.2	1376.0	80.0	9.2	948.0	172.0	172.0
3/10/2004	22:00:00	1.6	1272.0	51.0	6.5	836.0	131.0	131.0
...
4/4/2005	10:00:00	3.1	1314.0	-200.0	13.5	1101.0	472.0	472.0
4/4/2005	11:00:00	2.4	1162.0	-200.0	11.4	1027.0	353.0	353.0
4/4/2005	12:00:00	2.4	1142.0	-200.0	12.4	1062.0	293.0	293.0
4/4/2005	13:00:00	2.1	1002.0	-200.0	9.5	960.0	234.0	234.0
4/4/2005	14:00:00	2.2	1071.0	-200.0	11.9	1047.0	265.0	265.0

rows × 15 columns



```
In [24]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

```
In [25]: X = analysis_df.drop(['AH', 'Date', 'Time'],axis=1)
X
```

```
Out[25]:
```

	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	NOx(GT)	PT08.S3(NOx)	NC
0	2.6	1360.0	150.0	11.9	1046.0	166.0	1056.0	
1	2.0	1292.0	112.0	9.4	955.0	103.0	1174.0	
2	2.2	1402.0	88.0	9.0	939.0	131.0	1140.0	
3	2.2	1376.0	80.0	9.2	948.0	172.0	1092.0	
4	1.6	1272.0	51.0	6.5	836.0	131.0	1205.0	
...
9352	3.1	1314.0	-200.0	13.5	1101.0	472.0	538.0	
9353	2.4	1162.0	-200.0	11.4	1027.0	353.0	604.0	
9354	2.4	1142.0	-200.0	12.4	1062.0	293.0	603.0	
9355	2.1	1002.0	-200.0	9.5	960.0	234.0	702.0	
9356	2.2	1071.0	-200.0	11.9	1047.0	265.0	654.0	

9357 rows × 12 columns



```
In [26]: Y = analysis_df['AH']
Y
```

```
Out[26]: 0      1
1      1
2      1
3      1
4      1
..
9352   1
9353   1
9354   1
9355   1
9356   1
Name: AH, Length: 9357, dtype: int64
```

```
In [27]: X_train , X_test , Y_train , Y_test = train_test_split(X,Y,test_size=0.3)
```

```
In [28]: logreg = LogisticRegression()
logreg.fit(X_train,Y_train)
```

```
Out[28]: LogisticRegression
LogisticRegression()
```

```
In [29]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [30]: y_pred = logreg.predict(X_test)
```

```
In [31]: print(confusion_matrix(Y_test, y_pred))
```

```
[[ 114   0]
 [   0 2694]]
```

```
In [32]: print(classification_report(Y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	114
1	1.00	1.00	1.00	2694
accuracy			1.00	2808
macro avg	1.00	1.00	1.00	2808
weighted avg	1.00	1.00	1.00	2808