

Perform the following operations using 'python' Language on the Heart Disease dataset.

1. Data cleaning
2. Data integration
3. Data transformation
4. Data model Building

```
In [1]: import pandas as pd
import numpy as np
```

```
In [2]: df = pd.read_csv(r"C:\Users\yasha\Desktop\Ashish\sem 6\DSBDA\DSBDA Lab Dataset\heart.csv")
df
```

```
Out[2]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	
1020	59	1	1	140	221	0	1	164	1	0.0	2	0	2	
1021	60	1	0	125	258	0	0	141	1	2.8	1	1	3	
1022	47	1	0	110	275	0	0	118	1	1.0	1	1	2	
1023	50	0	0	110	254	0	0	159	0	0.0	2	0	2	
1024	54	1	0	120	188	0	1	113	0	1.4	1	1	3	

1025 rows × 14 columns



```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1025 non-null   int64
 1   sex         1025 non-null   int64
 2   cp          1025 non-null   int64
 3   trestbps    1025 non-null   int64
 4   chol        1025 non-null   int64
 5   fbs         1025 non-null   int64
 6   restecg     1025 non-null   int64
 7   thalach     1025 non-null   int64
 8   exang       1025 non-null   int64
 9   oldpeak     1025 non-null   float64
10   slope       1025 non-null   int64
11   ca          1025 non-null   int64
12   thal        1025 non-null   int64
13   target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

```
In [ ]: # data cleaning
```

```
In [4]: df.isnull().sum()
```

```
Out[4]: age         0
sex           0
cp            0
trestbps      0
chol          0
fbs           0
restecg       0
thalach       0
exang         0
oldpeak       0
slope         0
ca            0
thal          0
target        0
dtype: int64
```

```
In [7]: df.duplicated()
```

```
Out[7]: 0      False
        1      False
        2      False
        3      False
        4      False
        ...
       1020     True
       1021     True
       1022     True
       1023     True
       1024     True
Length: 1025, dtype: bool
```

```
In [10]: duplicate_rows = df[df.duplicated()]
         print("number of duplicate rows are :", duplicate_rows.shape)
```

```
number of duplicate rows are : (723, 14)
```

```
In [15]: df=df.drop_duplicates()
         duplicate_rows = df[df.duplicated()]
         print("number of duplicated rows are :",duplicate_rows.shape)
```

```
number of duplicated rows are : (0, 14)
```

```
In [ ]: #data integration
```

```
In [17]: df1=df[['age','sex','cp','ca']].loc[0:15]  
df1
```

```
Out[17]:
```

	age	sex	cp	ca
0	52	1	0	2
1	53	1	0	0
2	70	1	0	0
3	61	1	0	1
4	62	0	0	3
5	58	0	0	0
6	58	1	0	3
7	55	1	0	1
8	46	1	0	0
9	54	1	0	2
10	71	0	0	0
11	43	0	0	0
12	34	0	1	0
13	51	1	0	3
14	52	1	0	0

```
In [18]: df2=df[['age','sex','cp','ca']].loc[16:30]
df2
```

```
Out[18]:
```

	age	sex	cp	ca
16	51	0	2	1
17	54	1	0	1
18	50	0	1	0
19	58	1	2	0
20	60	1	2	0
21	67	0	0	2
22	45	1	0	0
23	63	0	2	0
24	42	0	2	0
25	61	0	0	0
26	44	1	2	0
27	58	0	1	2
28	56	1	2	1
29	55	0	0	0
30	44	1	0	0

```
In [19]: merge = pd.merge(df1,df2,on="age",how="inner")
merge
```

```
Out[19]:
```

	age	sex_x	cp_x	ca_x	sex_y	cp_y	ca_y
0	61	1	0	1	0	0	0
1	58	0	0	0	1	2	0
2	58	0	0	0	0	1	2
3	58	1	0	3	1	2	0
4	58	1	0	3	0	1	2
5	55	1	0	1	0	0	0
6	54	1	0	2	1	0	1
7	51	1	0	3	0	2	1

```
In [ ]: #data transformation
```

```
In [22]: df['target']=df['target'].apply(lambda x:1 if x>0 else 0)
df
```

```
Out[22]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
723	68	0	2	120	211	0	0	115	0	1.5	1	0	2	
733	44	0	2	108	141	0	1	175	0	0.6	1	0	2	
739	52	1	0	128	255	0	1	161	1	0.0	2	1	3	
843	59	1	3	160	273	0	0	125	0	0.0	2	0	2	
878	54	1	0	120	188	0	1	113	0	1.4	1	1	3	

302 rows × 14 columns



```
In [ ]: #error correction
```

```
In [24]: df =df.applymap(lambda x: df.mean() if x< 0 else x)
df
# if there is negative value it will replace it with mean of the data frame
```

```
Out[24]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
723	68	0	2	120	211	0	0	115	0	1.5	1	0	2	
733	44	0	2	108	141	0	1	175	0	0.6	1	0	2	
739	52	1	0	128	255	0	1	161	1	0.0	2	1	3	
843	59	1	3	160	273	0	0	125	0	0.0	2	0	2	
878	54	1	0	120	188	0	1	113	0	1.4	1	1	3	

302 rows × 14 columns



```
In [ ]: #model building
```

```
In [27]: from sklearn.model_selection import train_test_split
x = merge.drop(['age'], axis=1)
y = merge['age']
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_sta
```

```
In [29]: from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression()
logreg.fit(x_train,y_train)
```

```
Out[29]: LogisticRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [31]: from sklearn.metrics import classification_report, confusion_matrix
y_pred = logreg.predict(x_test)
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
[[0 0 0 0]
 [1 0 0 0]
 [0 0 1 0]
 [1 0 0 0]]
```

	precision	recall	f1-score	support
54	0.00	0.00	0.00	0
55	0.00	0.00	0.00	1
58	1.00	1.00	1.00	1
61	0.00	0.00	0.00	1
accuracy			0.33	3
macro avg	0.25	0.25	0.25	3
weighted avg	0.33	0.33	0.33	3

C:\Users\yasha\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\yasha\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1344: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\yasha\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\yasha\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1344: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\yasha\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\yasha\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1344: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))



